Advanced CUDA Webinar

Memory Optimizations

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
    - **Data transfers between host and device**
    - **Device memory optimizations**
        - **Measuring performance – effective bandwidth**
        - **Coalescing**
        - **Shared Memory**
        - **Textures**
- **Summary**

Optimize Algorithms for the GPU

- **Maximize independent parallelism**

- **Maximize arithmetic intensity (math/bandwidth)**

- **Sometimes it's better to recompute than to cache**
  - **GPU spends its transistors on ALUs, not memory**

- **Do more computation on the GPU to avoid costly data transfers**
  - **Even low parallelism computations can sometimes be faster than transferring back and forth to host**

Optimize Memory Access

- **Coalesced vs. Non-coalesced = order of magnitude**
  - **Global/Local device memory**

- **Optimize for spatial locality in cached texture memory**

- **In shared memory, avoid high-degree bank conflicts**

Take Advantage of Shared Memory

- Hundreds of times faster than global memory

- Threads can cooperate via shared memory

- Use one / a few threads to load / compute data shared by all threads

- Use it to avoid non-coalesced access
  - Stage loads and stores in shared memory to re-order non-coalesceable addressing

Use Parallelism Efficiently

- **Partition your computation to keep the GPU multiprocessors equally busy**
  - **Many threads, many thread blocks**

- **Keep resource usage low enough to support multiple active thread blocks per multiprocessor**
 - **Registers, shared memory**

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
  - **Data transfers between host and device**
  - **Device memory optimizations**
    - **Measuring performance – effective bandwidth**
    - **Coalescing**
    - **Shared Memory**
    - **Textures**
- **Summary**

10-Series Architecture

- **240 Scalar Processor (SP) cores execute kernel threads**
- **30 Streaming Multiprocessors (SMs) each contain**
  - **8 scalar processors**
  - **2 Special Function Units (SFUs)**
  - **1 double precision unit**
 - **Shared memory enables thread cooperation**

Multiprocessor

Scalar Processors

Double

Shared Memory

8

Execution Model

Software   Hardware

Thread

Scalar Processor

Threads are executed by scalar processors

Thread Block

Multiprocessor

Thread blocks are executed on multiprocessors

Thread blocks do not migrate

Several concurrent thread blocks can reside on one multiprocessor - limited by multiprocessor resources (shared memory and register file)

Grid

Device

A kernel is launched as a grid of thread blocks

Only one kernel can execute on a device at one time

Warps and Half Warps

Thread Block = Warps:
- 32 Threads
- 32 Threads
- 32 Threads

→ Multiprocessor

A thread block consists of 32-thread warps

A warp is executed physically in parallel (SIMD) on a multiprocessor

Half Warps: 16 | 16 ↔ DRAM (Global, Local) Device Memory

A half-warp of 16 threads can coordinate global memory accesses into a single transaction

Memory Architecture

Memory Architecture

Memory	Location	Cached	Access	Scope	Lifetime
Register	On-chip	N/A	R/W	One thread	Thread
Local	Off-chip	No	R/W	One thread	Thread
Shared	On-chip	N/A	R/W	All threads in a block	Block
Global	Off-chip	No	R/W	All threads + host	Application
Constant	Off-chip	Yes	R	All threads + host	Application
Texture	Off-chip	Yes	R	All threads + host	Application

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
  - **Data transfers between host and device**
  - **Device memory optimizations**
    - **Measuring performance – effective bandwidth**
    - **Coalescing**
    - **Shared Memory**
    - **Textures**
- **Summary**

Host-Device Data Transfers

- **Device to host memory bandwidth much lower than device to device bandwidth**
  - 8 GB/s peak (PCI-e x16 Gen 2) vs. 141 GB/s peak (GTX 280)

- **Minimize transfers**
  - Intermediate data can be allocated, operated on, and deallocated without ever copying them to host memory

- **Group transfers**
  - One large transfer much better than many small ones

Page-Locked Data Transfers

- `cudaMallocHost()` allows allocation of page-locked ("pinned") host memory

- Enables highest cudaMemcpy performance
  - 3.2 GB/s on PCI-e x16 Gen1
  - 5.2 GB/s on PCI-e x16 Gen2

- See the "bandwidthTest" CUDA SDK sample

- Use with caution!!
  - Allocating too much page-locked memory can reduce overall system performance
  - Test your systems and apps to learn their limits

Overlapping Data Transfers and Computation

- **Async and Stream APIs allow overlap of H2D or D2H data transfers with computation**
  - CPU computation can overlap data transfers on all CUDA capable devices
  - Kernel computation can overlap data transfers on devices with "Concurrent copy and execution" (roughly compute capability >= 1.1)

- **Stream = sequence of operations that execute in order on GPU**
  - Operations from different streams can be interleaved
  - Stream ID used as argument to async calls and kernel launches

Asynchronous Data Transfers

- **Asynchronous host-device memory copy returns control immediately to CPU**
  - `cudaMemcpyAsync(dst, src, size, dir, stream);`
  - requires *pinned* host memory (allocated with "`cudaMallocHost`")

- **Overlap CPU computation with data transfer**
  - **0** = default stream

```
cudaMemcpyAsync(a_d, a_h, size,
    cudaMemcpyHostToDevice, 0);
kernel<<<grid, block>>>(a_d);
cpuFunction();
```

overlapped

Overlapping kernel and data transfer

- **Requires:**
  - "**Concurrent copy and execute**"
    - `deviceOverlap` **field of a** `cudaDeviceProp` **variable**
  - **Kernel and transfer use different, *non-zero* streams**
    - **A CUDA call to stream-0 blocks until all previous calls complete and cannot be overlapped**

- **Example:**
  ```
  cudaStreamCreate(&stream1);
  cudaStreamCreate(&stream2);
  cudaMemcpyAsync(dst, src, size, dir, stream1);
  kernel<<<grid, block, 0, stream2>>>(…);
  ```
  overlapped

GPU/CPU Synchronization

- **Context based**
  - **cudaThreadSynchronize()**
    - **Blocks until all previously issued CUDA calls from a CPU thread complete**

- **Stream based**
  - **cudaStreamSynchronize(stream)**
    - **Blocks until all CUDA calls issued to given stream complete**

  - **cudaStreamQuery(stream)**
    - **Indicates whether stream is idle**
    - **Returns cudaSuccess, cudaErrorNotReady, ...**
    - **Does not block CPU thread**

GPU/CPU Synchronization

- **Stream based using events**
    - Events can be inserted into streams:

        `cudaEventRecord(event, stream)`
    - Event is recorded when GPU reaches it in a stream
        - Recorded = assigned a timestamp (GPU clocktick)
        - Useful for timing

    - `cudaEventSynchronize(event)`
        - Blocks until given event is recorded

    - `cudaEventQuery(event)`
        - Indicates whether event has recorded
        - Returns `cudaSuccess, cudaErrorNotReady, ...`
        - Does not block CPU thread

Zero copy

- **Access host memory directly from device code**
  - **Transfers implicitly performed as needed by device code**
  - **Introduced in CUDA 2.2**
  - **Check `canMapHostMemory` field of `cudaDeviceProp` variable**
- **All set-up is done on host using mapped memory**

```
cudaSetDeviceFlags(cudaDeviceMapHost);
...
cudaHostAlloc((void **)&a_h, nBytes,
    cudaHostAllocMapped);
cudaHostGetDevicePointer((void **)&a_d, (void *)a_h, 0);
for (i=0; i<N; i++) a_h[i] = i;
increment<<<grid, block>>>(a_d, N);
```

Zero copy considerations

- **Zero copy will always be a win for integrated devices that utilize CPU memory (you can check this using the `integrated` field in `cudaDeviceProp`)**
- **Zero copy will be faster if data is only read/written from/to global memory once, for example:**
  - **Copy input data to GPU memory**
  - **Run one kernel**
  - **Copy output data back to CPU memory**
- **Potentially easier and faster alternative to using `cudaMemcpyAsync`**
  - **For example, can both read and write CPU memory from within one kernel**
- **Note that current devices use pointers that are 32-bit so there is a limit of *4GB per context***

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
    - Data transfers between host and device
    - Device memory optimizations
        - Measuring performance – effective bandwidth
        - Coalescing
        - Shared Memory
        - Textures
- **Summary**

Theoretical Bandwidth

- **Device Bandwidth of GTX 280**

DDR

$$1107 * 10^6 * (512 / 8) * 2 / 1024^3 = 131.9 \text{ GB/s}$$

Memory clock (Hz)

Memory interface (bytes)

- **Specs report 141 GB/s**
  - **Use 10^9 B/GB conversion rather than 1024^3**
  - **Whichever you use, be consistent**

Effective Bandwidth

- **Effective Bandwidth (for copying array of N floats)**

  - $N * 4$ **B/element** $/ 1024^3 * 2 /$ **(time in secs)** $=$ **GB/s**

    Array size
    (bytes)

    Read and
    write

    B/GB
    (or 10^9)

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
    - **Data transfers between host and device**
    - **Device memory optimizations**
        - **Measuring performance – effective bandwidth**
        - **Coalescing**
        - **Shared Memory**
        - **Textures**
- **Summary**

Coalescing

- **Global memory access of 32, 64, or 128-bit words by a half-warp of threads can result in as few as one (or two) transaction(s) if certain access requirements are met**
- **Depends on compute capability**
  - **1.0 and 1.1 have stricter access requirements**
- **Float (32-bit) data example:**

32-byte segments

64-byte segments

128-byte segments

Global Memory

...

...

Half-warp of threads

Coalescing
Compute capability 1.0 and 1.1

- **K-th thread must access k-th word in the segment (or k-th word in 2 contiguous 128B segments for 128-bit words), not all threads need to participate**

Coalesces – 1 transaction

Out of sequence – 16 transactions

Misaligned – 16 transactions

Coalescing
Compute capability 1.2 and higher

- **Issues transactions for segments of 32B, 64B, and 128B**
- **Smaller transactions used to avoid wasted bandwidth**

1 transaction - 64B segment

2 transactions - 64B and 32B segments

1 transaction - 128B segment

Coalescing Examples

- **Effective bandwidth of small kernels that copy data**
  - **Effects of offset and stride on performance**

- **Two GPUs**
  - **GTX 280**
    - **Compute capability 1.3**
    - **Peak bandwidth of 141 GB/s**
  - **FX 5600**
    - **Compute capability 1.0**
    - **Peak bandwidth of 77 GB/s**

# Coalescing Examples

```
__global__ void offsetCopy(float *odata, float *idata,
                           int offset)
{
  int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;
  odata[xid] = idata[xid];
}
```


Copy with Offset

Coalescing Examples

```
__global__ void strideCopy(float *odata, float *idata,
                           int stride)
{
  int xid = (blockIdx.x*blockDim.x + threadIdx.x)*stride;
  odata[xid] = idata[xid];
}
```


Copy with Stride

Coalescing Examples

- **Strided memory access is inherent in many multidimensional problems**
  - **Stride is generally large (>> 18)**

However …

- **Strided access to global memory can be avoided using *shared memory***

Copy with Stride

GTX280
FX5600

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
    - **Data transfers between host and device**
    - **Device memory optimizations**
        - **Measuring performance – effective bandwidth**
        - **Coalescing**
        - **Shared Memory**
   - **Textures**
- **Summary**

Shared Memory

- ~**Hundred times faster than global memory**

- **Cache data to reduce global memory accesses**

- **Threads can cooperate via shared memory**

- **Use it to avoid non-coalesced access**
  - **Stage loads and stores in shared memory to re-order non-coalesceable addressing**

Shared Memory Architecture

- **Many threads accessing memory**
  - **Therefore, memory is divided into banks**
  - **Successive 32-bit words assigned to successive banks**

- **Each bank can service one address per cycle**
  - **A memory can service as many simultaneous accesses as it has banks**

- **Multiple simultaneous accesses to a bank result in a bank conflict**
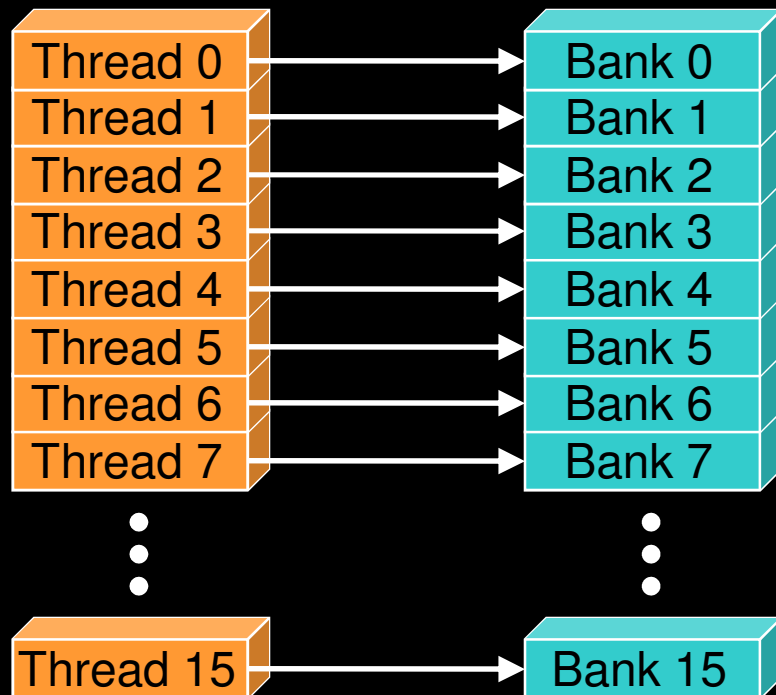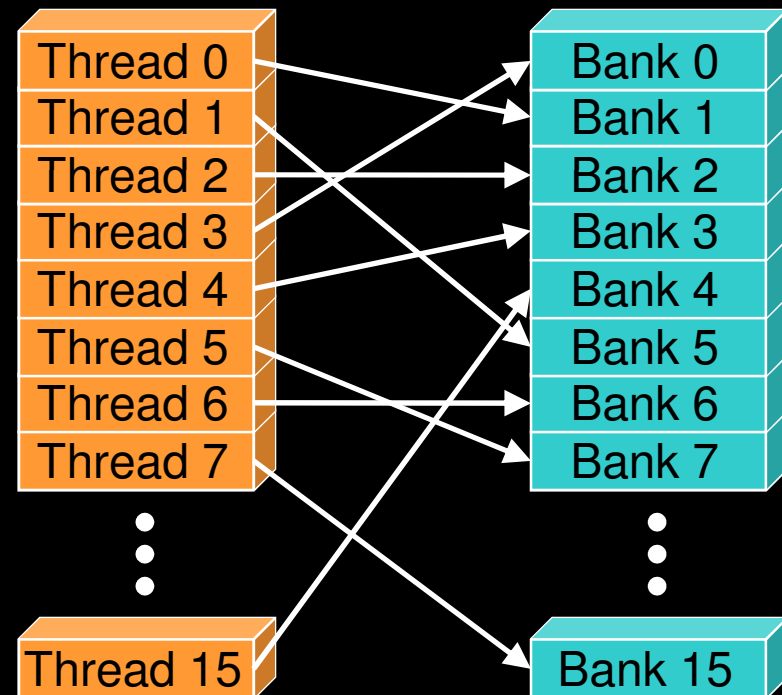 - **Conflicting accesses are serialized**

Bank 0
Bank 1
Bank 2
Bank 3
Bank 4
Bank 5
Bank 6
Bank 7

Bank 15

Bank Addressing Examples

- **No Bank Conflicts**
  - **Linear addressing stride == 1**

Thread 0	→	Bank 0
Thread 1	→	Bank 1
Thread 2	→	Bank 2
Thread 3	→	Bank 3
Thread 4	→	Bank 4
Thread 5	→	Bank 5
Thread 6	→	Bank 6
Thread 7	→	Bank 7
⋮		⋮
Thread 15	→	Bank 15

- **No Bank Conflicts**
  - **Random 1:1 Permutation**

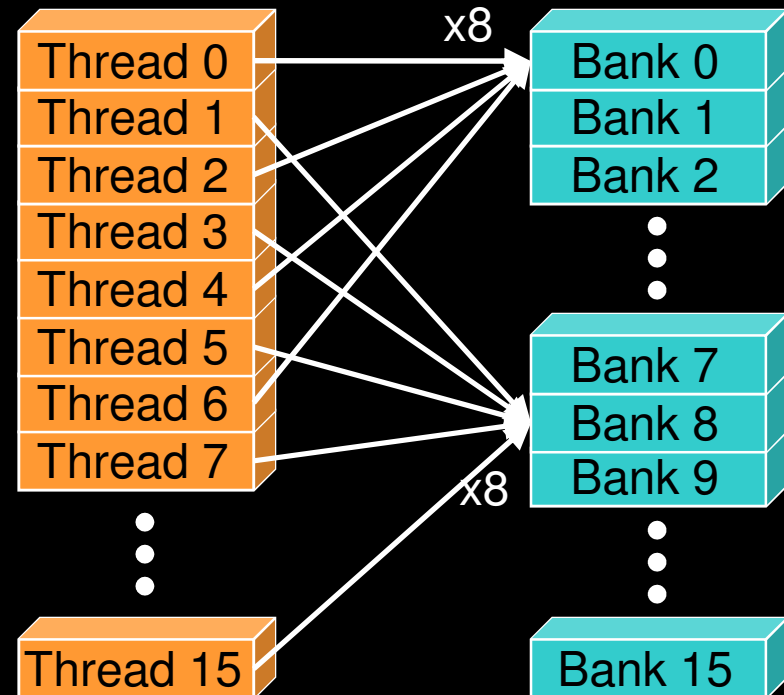| Thread 0 | Bank 0 |
Thread 1	Bank 1
Thread 2	Bank 2
Thread 3	Bank 3
Thread 4	Bank 4
Thread 5	Bank 5
Thread 6	Bank 6
Thread 7	Bank 7
⋮	⋮
Thread 15	Bank 15

Bank Addressing Examples

- **2-way Bank Conflicts**
  - **Linear addressing stride == 2**

- **8-way Bank Conflicts**
 - **Linear addressing stride == 8**

38

Shared memory bank conflicts

- **Shared memory is ~ as fast as registers if there are no bank conflicts**

- **warp_serialize profiler signal reflects conflicts**

- **The fast case:**
  - **If all threads of a half-warp access different banks, there is no bank conflict**
  - **If all threads of a half-warp read the identical address, there is no bank conflict (broadcast)**

- **The slow case:**
  - **Bank Conflict: multiple threads in the same half-warp access the same bank**
  - **Must serialize the accesses**
  - **Cost = max # of simultaneous accesses to a single bank**

Shared Memory Example: Transpose

- **Each thread block works on a tile of the matrix**
- **Naïve implementation exhibits strided access to global memory**

idata

odata

Elements transposed by a half-warp of threads

Naïve Transpose

- **Loads are coalesced, stores are not (strided by height)**

```
__global__ void transposeNaive(float *odata, float *idata,
                               int width, int height)
{
  int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;
  int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;

  int index_in  = xIndex + width * yIndex;
  int index_out = yIndex + height * xIndex;

  odata[index_out] = idata[index_in];
}
```

idata   odata

Coalescing through shared memory

- **Access columns of a tile in shared memory to write contiguous data to global memory**
- **Requires `__syncthreads()` since threads access data in shared memory stored by other threads**

idata   odata

tile

Elements transposed by a half-warp of threads

Coalescing through shared memory

```
__global__ void transposeCoalesced(float *odata, float *idata,
                                   int width, int height)
{
  __shared__ float tile[TILE_DIM][TILE_DIM];

  int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;
  int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;
  int index_in = xIndex + (yIndex)*width;

  xIndex = blockIdx.y * TILE_DIM + threadIdx.x;
  yIndex = blockIdx.x * TILE_DIM + threadIdx.y;
  int index_out = xIndex + (yIndex)*height;

  tile[threadIdx.y][threadIdx.x] = idata[index_in];

  __syncthreads();

  odata[index_out] = tile[threadIdx.x][threadIdx.y];
}
```

Bank Conflicts in Transpose

- **16x16 shared memory tile of floats**
  - **Data in columns are in the same bank**
  - **16-way bank conflict reading columns in tile**
- **Solution - pad shared memory array**
  - **`__shared__ float tile[TILE_DIM][TILE_DIM+1];`**
  - **Data in anti-diagonals are in same bank**

idata   odata

tile

Elements transposed by a half-warp of threads

Outline

- **Overview**
- **Hardware**
- **Memory Optimizations**
  - **Data transfers between host and device**
  - **Device memory optimizations**
    - **Measuring performance – effective bandwidth**
    - **Coalescing**
    - **Shared Memory**
    - **Textures**
- **Summary**

Textures in CUDA

- **Texture is an object for reading data**

- **Benefits:**
  - **Data is cached**
    - **Helpful when coalescing is a problem**
  - **Filtering**
    - **Linear / bilinear / trilinear interpolation**
    - **Dedicated hardware**
  - **Wrap modes (for "out-of-bounds" addresses)**
    - **Clamp to edge / repeat**
  - **Addressable in 1D, 2D, or 3D**
   - **Using integer or normalized coordinates**

Texture Addressing

0 1 2 3 4

(2.5, 0.5)
(1.0, 1.0)

Wrap

- Out-of-bounds coordinate is wrapped (modulo arithmetic)

0 1 2 3 4

(5.5, 1.5)

Clamp

- Out-of-bounds coordinate is replaced with the closest boundary

0 1 2 3 4

(5.5, 1.5)

CUDA Texture Types

- **Bound to linear memory**
  - Global memory address is bound to a texture
  - Only 1D
  - Integer addressing
  - No filtering, no addressing modes
- **Bound to CUDA arrays**
  - Block linear CUDA array is bound to a texture
  - 1D, 2D, or 3D
  - Float addressing (size-based or normalized)
  - Filtering
  - Addressing modes (clamping, repeat)
- **Bound to pitch linear (CUDA 2.2)**
  - Global memory address is bound to a texture
  - 2D
  - Float/integer addressing, filtering, and clamp/repeat addressing modes similar to CUDA arrays

CUDA Texturing Steps

- ## Host (CPU) code:
  - Allocate/obtain memory (global linear/pitch linear, or CUDA array)
  - Create a texture reference object
    - Currently must be at file-scope
  - Bind the texture reference to memory/array
  - When done:
    - Unbind the texture reference, free resources

- ## Device (kernel) code:
  - Fetch using texture reference
  - Linear memory textures: tex1Dfetch()
  - Array textures: tex1D() or tex2D() or tex3D()
  - Pitch linear textures: tex2D()

Texture Example

```
__global__ void
shiftCopy(float *odata,
          float *idata,
          int shift)
{
   int xid = blockIdx.x * blockDim.x
          + threadIdx.x;
   odata[xid] = idata[xid+shift];
}



texture <float> texRef;


__global__ void
textureShiftCopy(float *odata,
                 float *idata,
                 int shift)
{
   int xid = blockIdx.x * blockDim.x
          + threadIdx.x;
   odata[xid] = tex1Dfetch(texRef, xid+shift);
}
```


Copy with Shift
Using Global Memory and Textures

Summary

- **GPU hardware can achieve great performance on data-parallel computations if you follow a few simple guidelines:**
    - Use parallelism efficiently
    - Coalesce memory accesses if possible
    - Take advantage of shared memory
    - Explore other memory spaces
        - Texture
        - Constant
    - Reduce bank conflicts

Special CUDA Developer Offer on Tesla GPUs

- **50% off MSRP on Tesla C1060 GPUs**

- **Up to four per developer**

- **Act now, limited time offer**

- **Visit http://www.nvidia.com/object/webinar_promo**
  - *If you are outside of US or Canada, please contact an NVIDIA Tesla Preferred Provider in your country*